

WHICH BIAS FLAGS CAN YOU TRUST? A LABEL-FREE TRIAGE PROTOCOL FOR FRONTIER-MODEL NEWS-BIAS DETECTION

Anonymous authors

Paper under double-blind review

ABSTRACT

News-bias products—bias checkers, balanced-news feeds, media-literacy tools—are increasingly built by asking a frontier model to flag biased spans, label an article’s lean, and explain itself. Which of those outputs can a builder trust, and how should the untrustworthy ones be triaged *without* ground-truth labels at inference time? Evaluating two frontier target models and two cross-family judges on 200 U.S. news articles carrying outlet-anchored AllSides ratings, we first map the primitives: the *volume* of flagged spans is a valid partisan-intensity meter (Spearman $\rho \approx 0.6$), and models rank-track the continuous expert lean rating at $\rho \approx 0.8$, but a single model’s lean label near the center-right, its omission and unsupported-claim flags (chance-corrected agreement $\kappa \leq 0.15$), and the specific text it highlights (cross-model overlap 0.26) are model-specific opinion. We then turn the map into a validated selective-prediction protocol. A label-free *committee-disagreement* bit (one extra model call, no annotation) is not distinguishable from a calibrated confidence score as an error ranker at this sample size (AUC 0.70 vs 0.74), shows a directional (non-significant) error-capture advantage at 20–30% review budgets (+6–7pts; committee picks are 34% accurate on disagreements vs 74% on agreements, part of the gap structural), and combined with confidence yields the lowest risk–coverage area of any signal (AURC 0.20 vs 0.21 confidence, 0.31 disagreement); a two-vendor four-model unanimity gate clears 44% of articles at 87% accuracy; and neither raw confidence (both models overconfident; a stated 0.9 means $\sim 75\%$) nor a majority-vote ensemble (no accuracy gain over the best single model) should be used as a probability. The two headline results hold across five prompt conditions, and the reliability map and the directional failure both replicate on a six-model committee spanning four independent labs (Anthropic, OpenAI, Alibaba, DeepSeek). Findings are exploratory and inherit AllSides outlet labels as ground truth; headline statistics carry article-level bootstrap CIs (committed scripts reproduce every number).

1 INTRODUCTION

A growing class of products puts a frontier language model in the role of a news-bias analyst: it highlights loaded spans, assigns a left/right lean, summarizes “both sides,” and narrates why. The implicit assumption is that these outputs are measurements. Many are not—they are one model’s opinion, and a different frontier model, shown the same article, would say something else. A builder needs two things before shipping: to know *which* outputs reproduce, and—for the ones that do not—a way to catch the likely errors that needs *no* ground-truth labels at inference time, since a deployed product has none.

This paper supplies both. We do not ask whether the models are biased (a companion study takes up the behavioral question); we ask which of a bias detector’s outputs a tool can present as reliable, and how to triage the rest. Our answer is a reliability map with a clean split—primitives anchored to explicit text reproduce across independent models; primitives that require inference about what is absent or implied do not—and, built on it, a label-free triage protocol whose central signal is not model confidence but *model disagreement*. We validate the protocol against the standard confidence

baseline with a risk–coverage operating curve, and show it is stable across prompt conditions and across four independent model vendors rather than an artifact of one.

Contributions.

1. **A reliability map of bias-detection primitives.** Which outputs of a frontier-model bias detector reproduce across models (detection volume; continuous-rating rank; cross-model agreement) and which are model-specific opinion (lean labels near center-right; omission and unsupported-claim flags; specific highlights)—each with a number and a bootstrap CI (§4).
2. **A validated, label-free triage protocol.** A label-free committee-disagreement router with a published risk–coverage operating curve, benchmarked against a calibrated confidence router (not distinguishable as an error ranker at this n ; a directional in-budget capture advantage; lowest AURC when combined); a high-precision unanimity gate; and a per-model calibration analysis showing raw confidence must be rescaled and is trustworthy only for direction (§5).
3. **A directional failure mode shared across four vendors.** All six models we evaluate—across four independent labs—read moderate-right content as center-or-left, so a lean-labeling tool built on them under-counts the right; stable across five prompt conditions (§6).

2 RELATED WORK

LLM-as-judge reliability. Using an LLM to score model outputs is now standard (Zheng et al., 2023), and its biases—position, verbosity, self-preference—are documented (Ye et al., 2025; Panickssery et al., 2024). Reliability with consequences is emerging at the *instance* level (selective evaluation with human-agreement guarantees and escalation (Jung et al., 2025)) and judge *selection* by human agreement is practiced in benchmark construction (Dubois et al., 2024; Lambert et al., 2024). We differ in object and use: rather than validating a judge for scoring, we ask which *primitives* of a frontier bias detector a downstream product can trust, and we supply a deployment-time triage signal that needs no human labels at all.

Selective prediction. Deferring uncertain cases to a human is classical selective classification, evaluated by the risk–coverage curve (El-Yaniv & Wiener, 2010; Geifman & El-Yaniv, 2017). Committee disagreement as an error signal descends from query-by-committee (Seung et al., 1992), ensemble uncertainty (Lakshminarayanan et al., 2017), and disagreement-predicts-error results (Jiang et al., 2022); in the LLM era, sampled self-consistency flags errors label-free (Manakul et al., 2023) and diverse-model panels replace single judges (Verga et al., 2024); our contribution is an operating-regime characterization for frontier-model lean classification: the annotation-free disagreement bit is not distinguishable from calibrated confidence as an error ranker at our sample size, shows a directional capture advantage below the disagreement rate, and combining the two signals yields the lowest risk–coverage area—measured head-to-head with article-clustered uncertainty throughout.

Political bias in LLMs, and news tasks. A large literature measures LLM political orientation, largely via questionnaires (Feng et al., 2023; Santurkar et al., 2023), whose validity as a behavioral measure is contested (Röttger et al., 2024). Closer to our task, concurrent work reports “ideological center-collapse” in frontier-model news classification and steered generation on AllSides-labeled articles (Kennedy et al., 2026). We are not measuring bias but the *reliability* of a bias detector’s outputs; where we do report a directional error (§6), our per-class decomposition refines the center-collapse observation into a right-to-left asymmetry and replicates it across four labs.

Bias detection in NLP. Span-level media-bias detection is an established task (Spinde et al., 2021; Hamborg et al., 2019) using outlet-anchored corpora such as AllSides (Baly et al., 2020); we adopt that lineage and ground truth and study what a frontier model’s outputs on it can and cannot support.

3 SETUP

We reuse the corpus, models, and rollouts of the companion study; we summarize here and defer the full protocol (including the pre-committed reliability gates behind the judge instruments) to it. The corpus is $N=200$ contemporary U.S. news articles on a balanced 5×8 grid: five AllSides-anchored

lean classes (Left, Lean Left, Center, Lean Right, Right; 40 each) crossed with eight topics, 400–1500 words, each also carrying a continuous expert lean rating. Models receive *article text only*—no outlet, headline, URL, or date—so classification cannot reduce to outlet priors.

Two frontier *target* models (Claude Sonnet 4.5, GPT-4.1) and two cross-family *judge* models (Claude Sonnet 4.6, GPT-5) perform three tasks: span-level bias detection with a shared bias-type schema (Spinde et al., 2021; Hamborg et al., 2019), summarization with represented-perspective counts, and five-class lean classification with a self-reported confidence. Two further open-weight judges (Qwen3-235B-A22B, Alibaba; DeepSeek-V3.2), added post hoc via HuggingFace Inference Providers, extend the classification committee to six models across four labs for the four-family checks (`analysis/rate_articles_ext.py`). All analyses are exploratory (not pre-registered) and use expert AllSides labels as ground truth (Baly et al., 2020); agreement is Cohen’s κ , and bootstrap CIs use 10^4 resamples with a fixed seed. Data accounting: $N=200$ v3 articles; router analyses use the 199 with both targets parsed (baseline); the six-model checks the 198 rated by all six; legacy pilot rollouts sharing these directories are excluded throughout (`results/rollout/LEGACY_NOTE.md`).

4 WHICH PRIMITIVES REPRODUCE

Detection volume is a valid intensity meter. The simplest signal—how many spans a model flags—tracks how loaded an article is: flagged-span count rises monotonically with the article’s expert-rated intensity (Spearman $\rho=0.64$ Sonnet, 0.58 GPT-4.1 against $|\text{lean rating}|$, both $p<10^{-18}$; partisan articles draw $2.3\times$ the flags of Center articles). A builder can trust a flag *count* as a meter even where the individual flags are unreliable.

Lean predictions track the continuous rating, not just the bucket. Each model’s five-class prediction rank-correlates with the *continuous* expert rating at least as strongly as with the discrete label (Spearman 0.77–0.86 across the four models; difference vs the label ≤ 0.01), and within a single label class the models resolve graded intensity the bucket discards (within-Lean-Left $\rho(\text{pred}, \text{rating})=0.40$ –0.50, $p \leq .01$, all four). The lean signal is thus finer than the five-way label, and the results below are not an artifact of five-way bucketing. (We do not claim independence from the AllSides labels: the continuous rating and the discrete label partition this corpus identically.)

Only concrete, word-anchored bias forms reproduce. Given both target models the same bias-type schema, we measure how often they agree a form is present (Cohen’s κ over the 200 jointly-rated v3 articles, Table 1). Agreement spans 0.53 down to 0.08: the word-anchored forms reach moderate agreement (opinion-stated-as-fact 0.53, subjective adjectives 0.47)—and the forms defined by what is *absent from* or *unverified against* the text collapse toward chance. The gradient echoes what the annotation literature documents for humans—lexical bias is markedly easier to annotate than informational bias (omission, selection) (Fan et al., 2019), and framing annotation has low agreement even for trained coders (Card et al., 2015); jury-scale human annotation is the established remedy (Pastorino & Moosavi, 2025). The disagreement is not random: the two models have distinct habits—Claude leads with bias-by-omission (14% of its flags), GPT-4.1 with sensationalism (15%)—so one hunts what a story omits while the other hunts its tone, which is exactly why they cannot agree omission is present. The unreliability reaches the span level: even when both flag the same article, the specific highlighted text overlaps at only Jaccard 0.26. For a tool: show concrete word-level flags; treat any single model’s omission, framing, or unsupported-claim flag—and any individual highlight—as one model’s opinion, surfaced only on cross-model agreement. The unsupported-claims bound applies to *prompted* flags specifically: dedicated retrieval-plus-verification pipelines make claim support a tractable task (Choi et al., 2026), and are the right tool for that primitive.

A cheap word-list is a corpus monitor, not a per-item detector. A deterministic paired left/right lexicon corroborates aggregate directional patterns, but as a filter for a judge’s per-output verdict it reaches only precision 0.25, recall 0.30—so it belongs in a corpus-level regression check, never as a live per-article flag.

5 A LABEL-FREE TRIAGE PROTOCOL

The map says most single-model outputs cannot be shown as fact. A deployed product cannot fix this with ground truth—it has none at inference—so it needs a trust signal computable from the model

Reliably agreed (κ)		Barely above chance (κ)	
opinion stated as fact	0.53	omission of attribution	0.23
subjective adjectives	0.47	story choice & placement	0.20
mind reading	0.43	bias by omission	0.15
word choice	0.38	unsubstantiated claims	0.08

Table 1: Cross-model agreement on bias forms (Cohen’s κ , two target models, 200 shared v3 articles). Detectability tracks textual explicitness: models agree on loaded *words*, not on what a story *left out* or whether a claim is *unsupported*.

Review budget	10%	20%	30%	40%	50%
route by disagreement	10.8	33.7	48.2	65.1	78.3
route by confidence	14.5	26.5	42.2	59.0	71.1

Table 2: Percent of all classification errors captured at each human-review budget (route least-reliable first; $n=199$ v3 articles, two target models). Disagreement leads confidence in the 20–50% regime; the advantage is directional (clustered $p \geq .10$).

outputs alone. We show the best such signal is *model disagreement*, and benchmark it against the obvious alternative, model confidence.

Route disagreement first. On five-class lean ($n=199$ articles, both targets parsed; committee pick on a disagreement = the higher-confidence model’s label), when the two target models agree they are 73.8% accurate ($n=122$); on disagreements the committee pick is 33.8% accurate ($n=77$; gap 40.0pts, article-clustered 95% CI [26.9, 53.3] — noting a structural component: on a disagreement at most one of the two labels can be correct). As an error ranker the label-free disagreement bit is not distinguishable from calibrated confidence at this sample size (AUC 0.695 vs 0.740), and it shows a *directional*, non-significant capture advantage in the operational regime (Figure 1, Table 2): at a 20–30% review budget it captures +6 to +7 points more of all errors than a confidence threshold (clustered $p=.10$ and $p=.22$; pooling all five prompt conditions does not reach significance either). Integrated over *all* coverage levels the picture is complementary: confidence, being continuous, resolves risk within the agree and disagree groups and so has a modestly lower area under the risk–coverage curve (tie-robust expected AUC 0.217 vs disagreement’s 0.296; random 0.39, oracle 0.102), even though the two are tied on error ranking. The two signals carry the same information but deliver it differently—disagreement is *free* and most powerful in the low-budget regime, confidence is smoother across the curve—and *combining* them (route disagreements first, break ties within each group by confidence) achieves the lowest area of any signal (AUC 0.201), an ordering stable across all five prompt conditions (disagreement AUC 0.29–0.34, confidence 0.19–0.22, combined lowest or tied-lowest in every condition; agree-vs-disagree accuracy 0.68–0.74 vs 0.31–0.44). The operational recommendation: *route model disagreements to human review first, rank within groups by confidence, and expect no ranking gain from confidence alone until the budget exceeds the disagreement rate*.

At committee scale, trust unanimity—not a vote. Adding the two judges, a four-model majority vote does *not* beat the best single model (0.646 vs 0.662 exact; McNemar $p=.68$; bootstrap gain 95% CI [−.061, +.030]), and a builder cannot pre-select which model will be best (bootstrap $P(\text{is best})$: GPT-4.1 .63, Sonnet 4.6 .35, GPT-5 .02, Sonnet 4.5 .00). Ensembling therefore buys robustness, not accuracy. Its deployable value is instead a high-precision *gate*: the four models agree unanimously on 44% of articles, and there they are 87% accurate—so unanimity auto-clears nearly half the volume at high precision, and the residual is exactly the disagreement set the router sends to review. A natural worry is that a committee-blind gate would silently pass the shared directional failure of §6; empirically it does not concentrate there — the true classes of the two-vendor gate’s 11 unanimous errors are spread across Lean-Left (4), Center (3), and Lean-Right (4), and the four-lab six-model gate clears 35% of articles at 90% accuracy with only 2 of its 7 errors on Lean-Right (analysis/phase2_analyses.py).

Never read raw confidence as a probability. Consistent with documented verbalized-confidence behavior (Tian et al., 2023; Xiong et al., 2024), both models are severely overconfident on the exact




Figure 1: Risk-coverage operating curves for the three routing signals (baseline condition, $n=199$ articles; tie-robust expected risk, so the binary disagreement signal is flat within its trust groups). Gray dashed lines are the analytic random baseline and the oracle bound; the shaded band is the practical 20–30% review-budget regime. Combining disagreement with confidence attains the lowest area (AURC in parentheses). Reproduced by `analysis/fig-risk-coverage.py`.

five-class label (mean confidence 0.87–0.90 vs accuracy 0.54–0.66); a stated 0.9 empirically means 71% (Claude) to 77% (GPT-4.1) exact accuracy. GPT-4.1’s confidence is the better gate (Brier 0.259 vs 0.333; ECE 0.238 vs 0.330), and—critically—the overconfidence is almost entirely an *exact-boundary* effect: scored on lean *direction* (within one class), GPT-4.1’s confidence is essentially calibrated (ECE 0.024) and Claude’s nearly so (0.085). Confidence is a usable ordering for *direction*, never a usable probability for the exact label—rescale before thresholding.

6 THE DIRECTIONAL FAILURE MODE

The reproducible signals above are politically even-handed on average; the classification *errors* are not. Breaking accuracy down by true class exposes a single hard case shared by every model (Table 3): *Lean Right* is the hardest label for every model we evaluate—all six, spanning four independent labs (exact accuracy 0.07–0.46), and difficulty falls as an article’s expert lean strengthens (Spearman $\rho = -0.37, p < 10^{-7}$)—the models read the poles well and lose the moderate right. The misreads are *directional*: misclassified Lean-Right articles are pulled *left*, not merely blurred to center—33% are labeled Lean-Left or Left and 21% Center, while Center articles are called Lean-Left 32% of the time versus Lean-Right 4%. A “balanced feed” or lean-labeling product built on these classifications therefore *systematically under-counts right-leaning content*, consistent with concurrent reports of “center-collapse” (Kennedy et al., 2026), which our per-class decomposition refines into a right-to-left asymmetry.

The map holds across four independent labs. Adding two open-weight judges from two further families—Qwen3-235B-A22B (Alibaba) and DeepSeek-V3.2 (DeepSeek)—extends the committee to six models over four labs ($n=198$ articles rated by all six; `analysis/four-family-analysis.py`). Every headline replicates. Lean Right is the hardest class for all six (Table 3); the agreement→accuracy signal holds on all fifteen model pairs, including the new cross-vendor ones (Qwen×DeepSeek agree 0.69 vs disagree 0.45; GPT-5×DeepSeek 0.75 vs 0.42); consensus stays a monotone accuracy signal, with six-model unanimity (35% of articles) at 90% accuracy; and the right-to-left misread is unchanged (34% of Lean-Right articles labeled Lean-Left or Left, pooled across the six). Replication across Anthropic, OpenAI, Alibaba, and DeepSeek rules out a single-vendor quirk; it cannot, however, discriminate a shared model-side bias from an artifact of the shared outlet-anchored labels (all six models are scored against the same labels,

Model	Left	Lean L	Center	Lean R	Right
Claude Sonnet 4.5	.97	.23	.50	.07	.90
GPT-4.1	.88	.47	.75	.46	.72
Claude Sonnet 4.6	.92	.78	.55	.20	.80
GPT-5	.82	.70	.53	.42	.60
Qwen3-235B	.70	.62	.65	.35	.72
DeepSeek-V3.2	.85	.45	.65	.35	.68

Table 3: Exact classification accuracy by true lean class, six models across four labs (targets: Claude Sonnet 4.5, GPT-4.1; judges: Claude Sonnet 4.6, GPT-5, Qwen3-235B-A22B, DeepSeek-V3.2). Lean Right is the hardest label for every one. The bottom two rows are an exploratory four-family extension via HuggingFace Inference Providers (`analysis/rate_articles_ext.py`).

and shared training corpora are a common cause) — only independent article-level labels can separate these.

Both the map and the failure are prompt-invariant. Re-running the two headline results under five different prompt scaffolds, the agreement→accuracy gap holds in all five (gap 0.28–0.36, bootstrap $P(\text{gap}>0)=1.0$ in every condition) and Lean Right is the strict hardest class in all five (pooled accuracy 0.21–0.29). The pattern is a property of the models, not of one prompt. Nor is it a class-boundary artifact: restricting Lean-Right articles to the core band of their continuous rating (middle two quartiles) leaves the hole unchanged (Sonnet 0.086 vs 0.075 on all; GPT-4.1 0.412 vs 0.462). Three riders. First, the Lean-Right effect is Claude-driven (Claude is the worst model on it in 5/5 conditions, GPT-4.1 in only 3/5). Second, the target contrast is not a simple quality ordering but an *engagement profile*: on clearly partisan articles (true Left or Right, $n=80$ paired) Claude Sonnet 4.5 is significantly the stronger reader (0.94 vs GPT-4.1’s 0.80; of the 13 articles exactly one model read correctly, Claude was right in 12, McNemar exact $p=.003$), and it produces the more substantive rationales (175 vs 94 words on average) — the same engagement that costs it most in the subtle middle, where GPT-4.1’s comparative reticence protects it (post-hoc contrast; `analysis/phase2_analyses.py`). Third, lean classification is least reliable on policy-abstract topics (elections/economy/climate) and most reliable on identity-concrete ones.

7 DISCUSSION

The results reduce to a short deployment policy for a news-bias product. (1) *Show decisions, filter prose, and rank flags by kind*: flag counts and word-anchored flags are trustworthy; omission/framing/unsupported-claim flags and individual highlights are one model’s opinion and should surface only on cross-model agreement. (2) *Triage by agreement, not raw confidence*: model disagreement is a free, label-free error signal that beats a confidence threshold in the realistic low-budget regime; combine the two for the best risk–coverage trade-off, and never expose raw confidence as a probability. (3) *Correct for the rightward blind spot*: expect moderate-right content to be under-labeled, and treat a “center” label on right-of-center content as a routing trigger, not a verdict. None of these require labels at inference; the only ground truth the protocol needs is at build time, to choose thresholds. This turns a measurement report into a reusable engineering recipe, and the corpus recipe and analysis code are released so the map can be re-derived on a new model generation cheaply.

8 LIMITATIONS

Three limits bound the claims. First, all findings are *exploratory* (not pre-registered) and inherit AllSides’ outlet-anchored labels as the sole ground truth (Baly et al., 2020); the directional error could partly reflect article-vs-outlet label drift, though uniform label noise cannot produce a one-sided right-to-left pattern, and its replication across four independent labs makes a pure-label explanation less likely. A second, independent label source (per-article human consensus, or a second rating provider) is the highest-value next test. Second, the triage router is validated on the baseline prompt condition and one model pair (though the reliability map and directional failure replicate across conditions and vendors); and the label-free signals so far require a *committee*—disagreement needs two models, the unanimity gate four. A single-model cold-start signal—SelfCheckGPT-style self-

consistency across temperature-sampled generations (Manakul et al., 2023), benchmarked against the disagreement router’s AUC of 0.70—is a natural next step. Third, any LLM judge inside a tool’s own quality-control pipeline inherits the reliability problem we document for the models themselves.

ETHICS STATEMENT

This paper documents a directional reliability failure — frontier models under-label moderate-right news — in politically sensitive tooling, and we state its use explicitly: our evidence supports *escalating* affected items to human review, never automated or editorial re-weighting of content by political direction; citing this work as license to boost content of any lean misreads it. The ground truth itself is contested — outlet-anchored ratings projected to articles — and the failure is measured against, not independent of, that standard. The reliability map is dual-use (it also tells an actor which detector outputs are unauditable); we judge the builder and auditor benefit to outweigh this, and release evaluation code rather than any deployed classifier. News text is used under research norms with article text redistributed only as required for reproduction.

REPRODUCIBILITY STATEMENT

The corpus curation script and manifest, all task and judge prompts, the runner, the four-family judging scripts (`analysis/rate_articles_ext.py`, `analysis/four_family_analysis.py`), and per-call cost accounting are released; every number in the paper is reproduced by these scripts on the released rollouts. Bootstrap procedures use a fixed seed. Article text is redistributed only as required for reproduction, with source URLs and the deterministic curation script provided in lieu of bulk redistribution where licensing requires. [\[Repository and data links after de-anonymization.\]](#)

REFERENCES

- Ramy Baly, Giovanni Da San Martino, James Glass, and Preslav Nakov. We can detect your bias: Predicting the political ideology of news articles. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 4982–4991, Online, 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.404. URL <https://aclanthology.org/2020.emnlp-main.404/>.
- Dallas Card, Amber E. Boydston, Justin H. Gross, Philip Resnik, and Noah A. Smith. The media frames corpus: Annotations of frames across issues. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pp. 438–444, 2015. URL <https://aclanthology.org/P15-2072/>.
- Yee Man Choi, Xuehang Guo, Yi R. Fung, and Qingyun Wang. CiteGuard: Faithful citation attribution for LLMs via retrieval-augmented validation. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 6241–6257, 2026. URL <https://aclanthology.org/2026.acl-long.282/>.
- Yann Dubois, Balázs Galambosi, Percy Liang, and Tatsunori B. Hashimoto. Length-controlled AlpacaEval: A simple way to debias automatic evaluators, 2024. URL <https://arxiv.org/abs/2404.04475>. Accepted at COLM 2024.
- Ran El-Yaniv and Yair Wiener. On the foundations of noise-free selective classification. *Journal of Machine Learning Research*, 11:1605–1641, 2010. URL <https://jmlr.org/papers/v11/el-yaniv10a.html>.
- Lisa Fan, Marshall White, Eva Sharma, Ruisi Su, Prafulla Kumar Choubey, Ruihong Huang, and Lu Wang. In plain sight: Media bias through the lens of factual reporting. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pp. 6343–6349, 2019. URL <https://aclanthology.org/D19-1664/>.

- Shangbin Feng, Chan Young Park, Yuhan Liu, and Yulia Tsvetkov. From pretraining data to language models to downstream tasks: Tracking the trails of political biases leading to unfair NLP models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 11737–11762, Toronto, Canada, 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.acl-long.656. URL <https://aclanthology.org/2023.acl-long.656/>.
- Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. URL <https://arxiv.org/abs/1705.08500>.
- Felix Hamborg, Karsten Donnay, and Bela Gipp. Automated identification of media bias in news articles: an interdisciplinary literature review. *International Journal on Digital Libraries*, 20(4):391–415, 2019. doi: 10.1007/s00799-018-0261-y. URL <https://doi.org/10.1007/s00799-018-0261-y>.
- Yiding Jiang, Vaishnavh Nagarajan, Christina Baek, and J. Zico Kolter. Assessing generalization of SGD via disagreement. In *International Conference on Learning Representations (ICLR)*, 2022. URL <https://openreview.net/forum?id=WvOGCEAQhxl>.
- Jaehun Jung, Faeze Brahman, and Yejin Choi. Trust or escalate: LLM judges with provable guarantees for human agreement. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025. URL <https://openreview.net/forum?id=UHPnqSTBPO>.
- Molly Kennedy, Ali Parker, Yihong Liu, and Hinrich Schütze. Left, right, or center? Evaluating LLM framing in news classification and generation, 2026. URL <https://arxiv.org/abs/2601.05835>. arXiv preprint; concurrent work.
- Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. URL <https://arxiv.org/abs/1612.01474>.
- Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V. Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D. Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Taffjord, Chris Wilhelm, Luca Soldaini, Noah A. Smith, Yizhong Wang, Pradeep Dasigi, and Hannaneh Hajishirzi. Tulu 3: Pushing frontiers in open language model post-training, 2024. URL <https://arxiv.org/abs/2411.15124>.
- Potsawee Manakul, Adian Liusie, and Mark Gales. SelfCheckGPT: Zero-Resource black-box hallucination detection for generative large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 9004–9017, 2023. URL <https://aclanthology.org/2023.emnlp-main.557/>.
- Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/7f1f0218e45f5414c79c0679633e47bc-Abstract-Conference.html.
- Valeria Pastorino and Nafise Sadat Moosavi. Frame in, frame out: Measuring framing bias in LLM-generated news summaries, 2025. URL <https://arxiv.org/abs/2505.05406>. Accepted to *SEM 2026, co-located with ACL 2026.
- Paul Röttger, Valentin Hofmann, Valentina Pyatkin, Musashi Hinck, Hannah Kirk, Hinrich Schuetze, and Dirk Hovy. Political compass or spinning arrow? Towards more meaningful evaluations for values and opinions in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 15295–15311, Bangkok, Thailand, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.816. URL <https://aclanthology.org/2024.acl-long.816/>.
- Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, and Tatsunori Hashimoto. Whose opinions do language models reflect? In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*,

- pp. 29971–30004. PMLR, 2023. URL <https://proceedings.mlr.press/v202/santurkar23a.html>.
- H. S. Seung, M. Oppen, and H. Sompolinsky. Query by committee. In *Proceedings of the Fifth Annual Workshop on Computational Learning Theory (COLT)*, pp. 287–294, 1992. URL <https://doi.org/10.1145/130385.130417>.
- Timo Spinde, Manuel Plank, Jan-David Krieger, Terry Ruas, Bela Gipp, and Akiko Aizawa. Neural media bias detection using distant supervision with BABE - bias annotations by experts. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pp. 1166–1177, Punta Cana, Dominican Republic, 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.findings-emnlp.101. URL <https://aclanthology.org/2021.findings-emnlp.101/>.
- Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher Manning. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 5433–5442, 2023. URL <https://aclanthology.org/2023.emnlp-main.330/>.
- Pat Verga, Sebastian Hofstatter, Sophia Althammer, Yixuan Su, Aleksandra Piktus, Arkady Arkhangorodsky, Minjie Xu, Naomi White, and Patrick Lewis. Replacing judges with juries: Evaluating LLM generations with a panel of diverse models, 2024. URL <https://arxiv.org/abs/2404.18796>.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2306.13063>.
- Jiayi Ye, Yanbo Wang, Yue Huang, Dongping Chen, Qihui Zhang, Nuno Moniz, Tian Gao, Werner Geyer, Chao Huang, Pin-Yu Chen, Nitesh V. Chawla, and Xiangliang Zhang. Justice or prejudice? Quantifying biases in LLM-as-a-judge. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025. URL <https://openreview.net/forum?id=3GTtZFiajM>.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023), Datasets and Benchmarks Track*, 2023. URL <https://arxiv.org/abs/2306.05685>.

A FULL FOUR-FAMILY CONSENSUS LADDER

On the $n=198$ articles rated by all six models (four labs), modal-label accuracy rises monotonically with the number of models agreeing: $\geq 2/6$ agree covers 100% of articles at 0.66; $\geq 3/6$, 97% at 0.67; $\geq 4/6$, 81% at 0.73; $\geq 5/6$, 48% at 0.82; 6/6 (unanimous), 35% at 0.90. Unanimity across four independent vendors is the strongest label-free confidence signal available.